


Mathematical Software Programming (02635)

Module 7 – Fall 2016

Instructor: Bernd Dammann

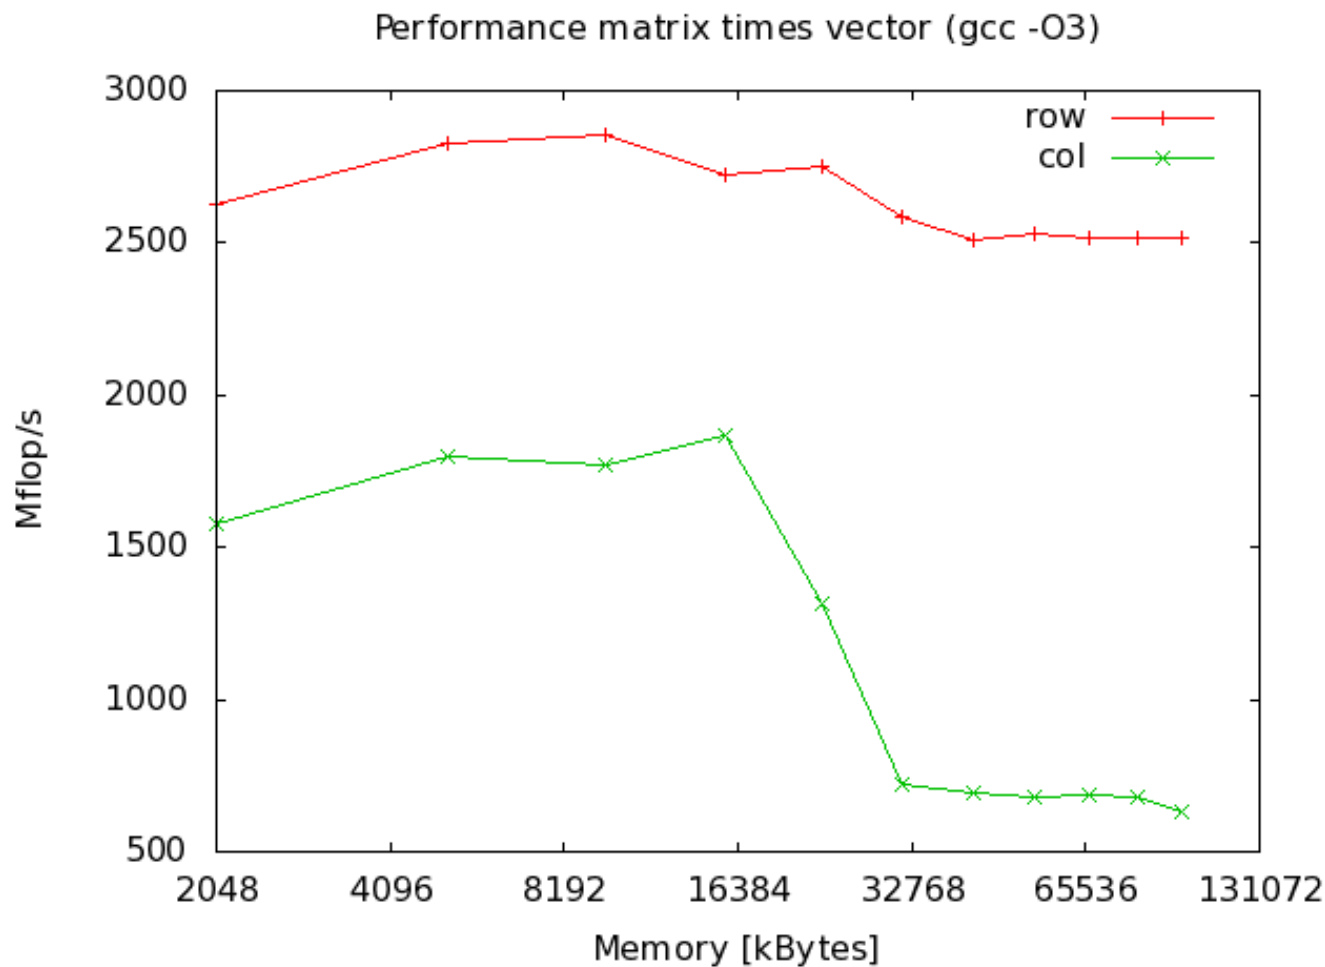
Re-cap from last week

Matrix times vector (runtimes):




Re-cap from last week

Matrix times vector (performance):



Re-cap from last week

Accessing vector elements in C:




size of a
cache line


Legend:



vector element



cache miss




cache hit







memory access

Re-cap from last week

Accessing 2d arrays in C – row wise:



Legend:

-  vector element
-  cache miss
-  cache hit
-  memory access

Re-cap from last week

Accessing 2d arrays in C – column wise:



Today's topics

- ❑ Parallelism – what is that?
- ❑ Parallel execution models
- ❑ Parallel speed-up – what is that, and what can we expect?
- ❑ Exploiting parallelism with OpenMP

Today's goal

- ❑ Basic understanding of parallel computations.
- ❑ Implement a parallel version of last week's code matrix times vector, using OpenMP.

What is Parallelization?

An attempt of a definition:

“Something” is parallel, if there is a certain level of independence in the order of operations

“Something” can be:

- ▶ A collection of program statements
- ▶ An algorithm
- ▶ A part of your program
- ▶ The problem you are trying to solve




Parallelism is everywhere

In today's computer installations one has many levels of parallelism:

- ❑ Instruction level (ILP)
- ❑ Chip level (multi-core, multi-threading)
- ❑ System level (multi-socket, i.e. multi-CPU)
- ❑ accelerators: GPU, Intel Xeon Phi, FPGA
- ❑ Cluster: “network of compute nodes”
- ❑ ...

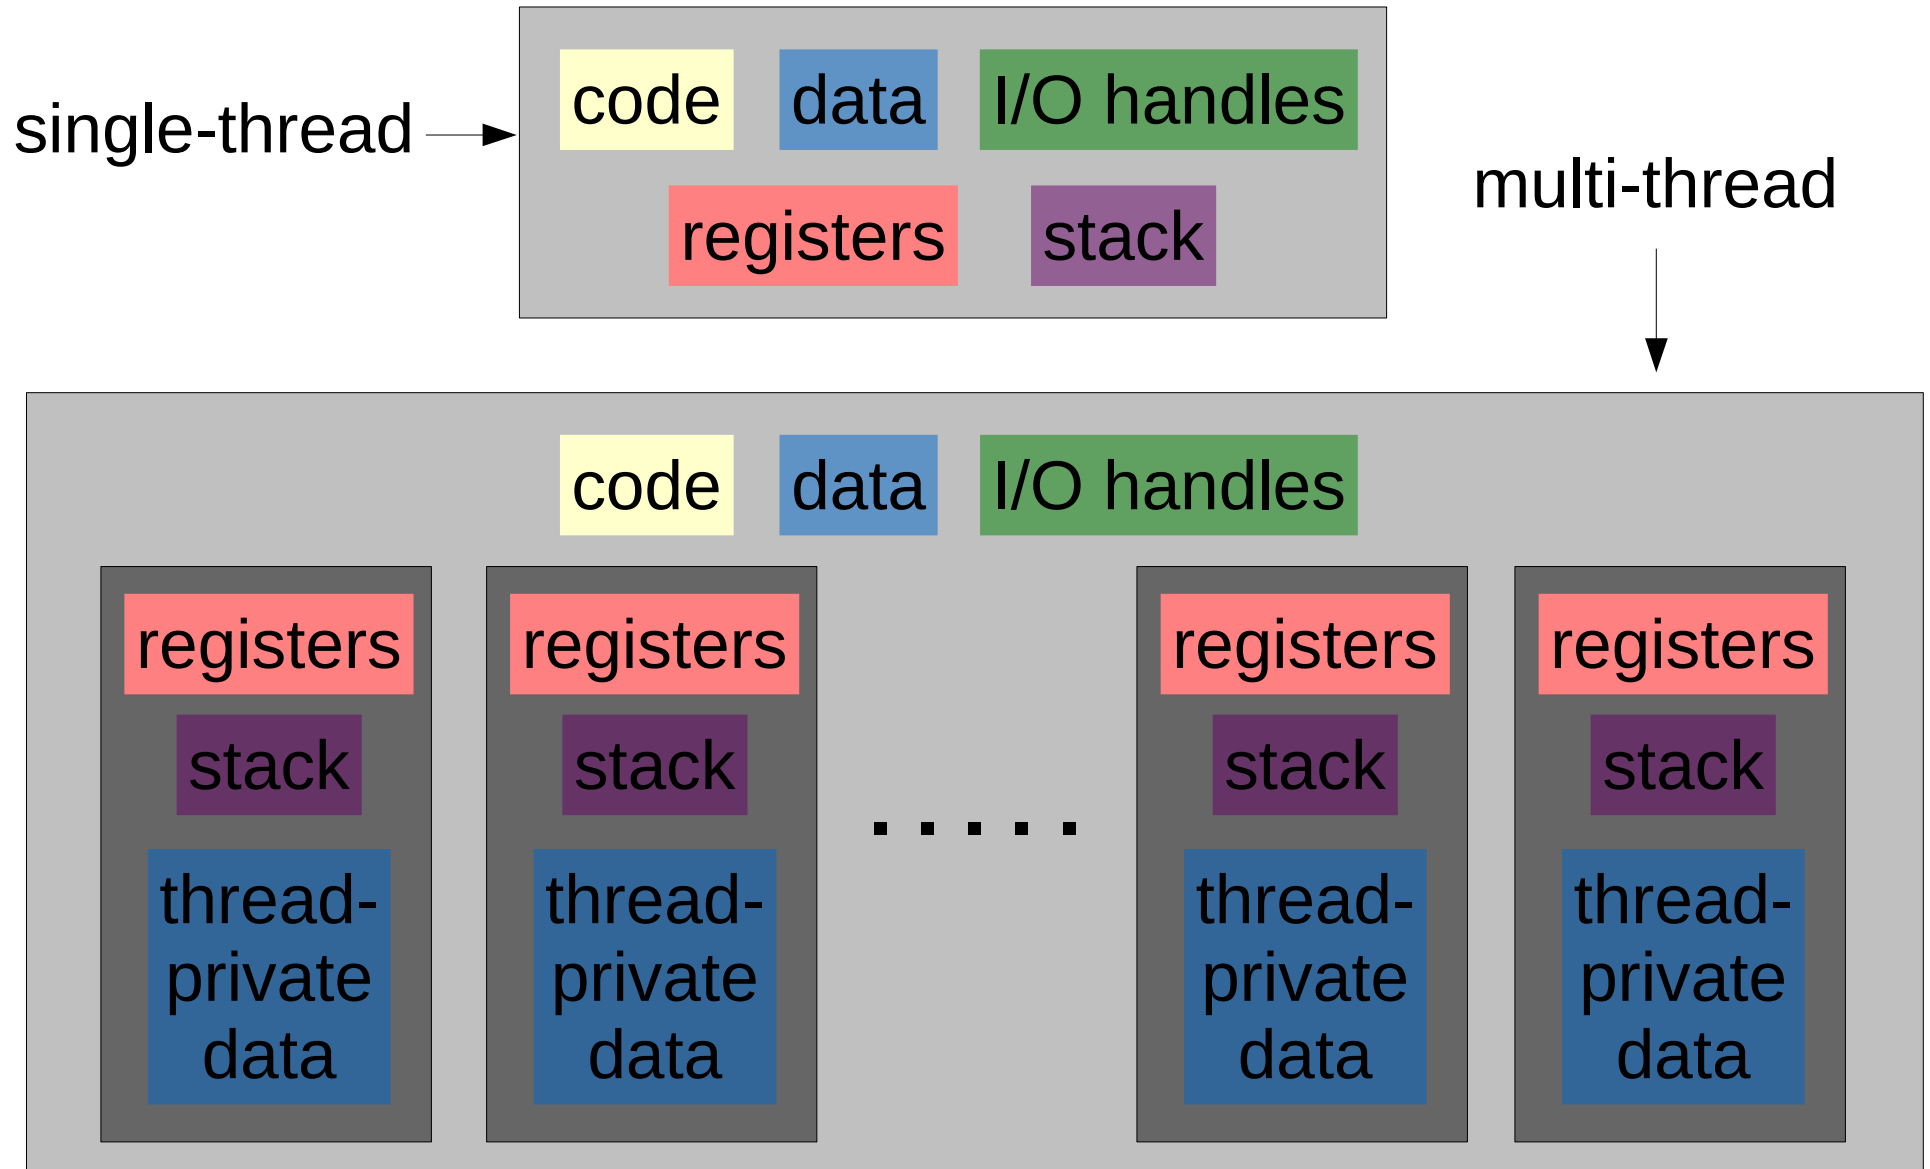
A typical multi-core setup



a 2-socket, 12-core, 24-(hyper-)threads server \Rightarrow 24 logical CPUs


Note: we do not use hyperthreading in our setup!

Single- vs. multi-threaded



What is a thread?

- ❑ Loosely said, a thread consists of a series of instructions with its own program counter (“PC”) and state
- ❑ A parallel program will execute threads in parallel
- ❑ These threads are then scheduled onto processors by the OS



Parallel execution models

❑ Multi-threaded:

- ❑ one process
- ❑ multiple threads
- ❑ “communication” (implicit) via shared-memory (shm)
- ❑ limited to one node (computer)

❑ Multi-process:

- ❑ multiple processes (usually single threaded)
- ❑ communication via interconnect (network or shm)
- ❑ can run on “any” number of nodes

❑ Hybrid: multiple multi-threaded processes

Timings in parallel programs


- ❑ So far, we have used `clock()` to time the speed of our programs, i.e. the CPU time
- ❑ In parallel programs:
 - ❑ the CPU time will very likely go up (parallel overhead)
 - ❑ `clock()` measures the accumulated time of all threads(!)
 - ❑ we need another measure: wallclock time, i.e. the time the user has to wait to get the result
- ❑ All parallel programming models provide a function to get the wallclock time.
- ❑ On the next slides: wall-time = wallclock time

Parallelism: speed-up

- ❑ What is this “speed-up”?
 - ❑ $S(p) = (\text{wall-time on 1 core}) / (\text{wall-time on } p \text{ cores})$
 $= T(1) / T(p)$
- ❑ ideal case: linear speed-up, e.g. $S(p) = p$
 - ❑ but: the world is not ideal
 - ❑ parallel overhead: extra instructions, communication, synchronization, etc
 - ❑ not all parts of your code can run in parallel – there will always be sequential code
 - ❑ in general: wall-time goes down – but CPU time goes up!

Parallelism: speed-up

- let f be the **parallel fraction** of your code, and $(1-f)$ the **sequential part**, e.g. $f = 0.5$



$T(1)$


- What is the max. speed-up, if we had an infinite number of cores ($p = \infty$), and no communication costs, etc?




$T(p=\infty)$

$$S = T(1) / T(p=\infty) < 2$$

Parallelism: Amdahl's law



Parallelism: Amdahl's law in practice





**Commercial
Break**

High-Performance Computing

- ❑ Got interested? Want to learn more?
- ❑ 02614 – High-Performance Computing
 - ❑ 3 weeks course in January
 - ❑ week 1: Serial tuning
 - ❑ week 2: Parallel programming with OpenMP
 - ❑ week 3: Scientific computing on GPUs
- ❑ 02616 – Large-scale modeling (in spring)
 - ❑ go beyond one node (MPI, etc)

Exploiting parallelism using OpenMP

What is OpenMP?

From openmp.org:

“The OpenMP API supports multi-platform shared-memory parallel programming in C/C++ and Fortran. The OpenMP API defines a portable, scalable model with a simple and flexible interface for developing parallel applications on platforms from the desktop to the supercomputer.”

- ❑ OpenMP is a “kind of add-on” to C/C++, Fortran
- ❑ it is not a programming language
- ❑ it requires a compiler that supports OpenMP

OpenMP components

- ❑ Directives
 - ❑ in your source code
 - ❑ e.g. parallel for-loop
- ❑ Environment variables
 - ❑ control program behaviour at runtime
 - ❑ e.g. number of threads to be used
- ❑ Runtime library
 - ❑ support functions
 - ❑ e.g. wallclock timer, etc

OpenMP: Hello world

OpenMP version of “Hello world”:

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    #pragma omp parallel
    {
        printf("Hello parallel world!\n");
    } /* end parallel */
    return(0);
}
```

OpenMP: Hello world

Compile and run ...

```
$ cc -o hello_omp hello_omp.c
```

```
$ ./hello_omp  
Hello parallel world!
```

```
$ OMP_NUM_THREADS=2 ./hello_omp  
Hello parallel world!
```

OpenMP: Hello world

Compile with OpenMP enabled – and run ...

```
$ cc -fopenmp -o hello_omp hello_omp.c
```

```
$ ./hello_omp  
Hello parallel world!
```

```
$ OMP_NUM_THREADS=2 ./hello_omp  
Hello parallel world!  
Hello parallel world!
```

OpenMP: Hello world v2

```
#include <stdio.h>
#ifdef _OPENMP
#include <omp.h>
#endif

int main(int argc, char *argv[]) {
    int t_id = 0;
    #pragma omp parallel private(t_id)
    {
        #ifdef _OPENMP
        t_id = omp_get_thread_num();
        #endif
        printf("Hello world from %d!\n", t_id);
    } /* end parallel */
    return(0);
}
```

OpenMP: Hello World v2

```
$ ./hello_omp2  
Hello world from 0!
```

```
$ OMP_NUM_THREADS=4 ./hello_omp2  
Hello world from 0!  
Hello world from 3!  
Hello world from 1!  
Hello world from 2!
```

- ❑ Note: The order of execution will be different from run to run!
- ❑ The default no. of threads depends on the OpenMP implementation

OpenMP: Parallel for-loop

Work-sharing – Loop parallelism:

- ❑ OpenMP implements parallel do/for-loops only!

```
int i;  
float a[N], b[N], c[N];  
  
for (i=0; i < N; i++)  
    a[i] = b[i] = i * 1.0;
```

```
#pragma omp parallel shared(a,b,c) private(i)  
{  
    #pragma omp for  
    for (i=0; i < N; i++)  
        c[i] = a[i] + b[i];  
} /* end of parallel region */
```

*for has to follow the
pragma – no {... }!*

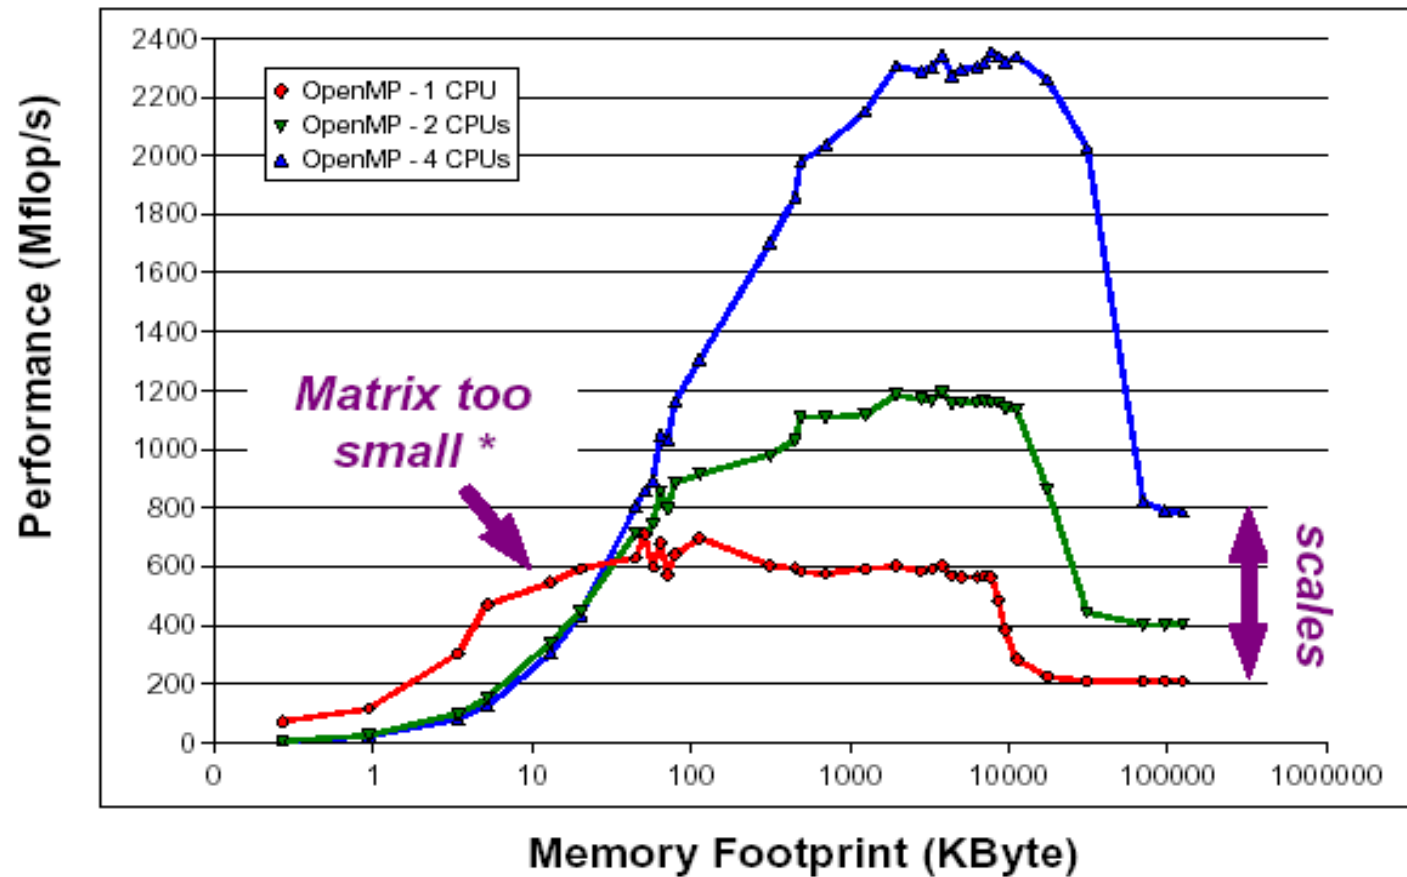
OpenMP: Parallel for-loop

Work-sharing – Loop parallelism:

- Another version: combined “parallel for”

```
int i;  
float a[N], b[N], c[N];  
  
for (i=0; i < N; i++)  
    a[i] = b[i] = i * 1.0;  
  
#pragma omp parallel for shared(a,b,c) \  
                        private(i)  
for (i=0; i < N; i++)  
    c[i] = a[i] + b[i];
```

OpenMP: Matrix times vector



SunFire 6800
UltraSPARC III Cu @ 900 MHz
8 MB L2-cache

courtesy: Ruud van der Pas, Oracle

Summary

Summary: Parallelism

- ❑ Parallel execution can speed up your code
- ❑ Wallclock time goes down – but the CPU time goes usually up (more resources, parallel overhead!)
- ❑ Don't expect magic ...
 - ❑ remember Amdahl's law!
 - ❑ is your problem too small?
 - ❑ don't use too many threads!
- ❑ Always check your results – compare to serial version!

Today's exercises

- ❑ Make your first parallel steps:
 - ❑ implement the “Hello World” example from the lecture
 - ❑ this should help you to understand how OpenMP works with your compiler
- ❑ Make parallel versions of last week's examples
 - ❑ row-wise version
 - ❑ column-wise version