

# Mathematical Software Programming (02635)

Module 9 — Fall 2016

Instructor: Martin S. Andersen

# Announcement

## Assignment 2

- ▶ Posted on CampusNet
- ▶ Due on November 23, 2016 (hand-in through CampusNet)
- ▶ Please post your questions on **Piazza**

### Learning objectives

- ▶ call external (third party) programs and libraries (\*)
- ▶ design, implement, and document a program that solves a mathematical problem
- ▶ debug and **test** mathematical software

(\*) Only assessed through this assignment (and **not** through final exam)

# This week

## Topics

- ▶ External libraries and basic linear algebra

## Learning objectives

- ▶ Call external (third party) programs and libraries

# Basic Linear Algebra Subroutines (BLAS)

- ▶ building blocks for linear algebra (*de facto* standard)
- ▶ started as a FORTRAN library (late 1970s)
- ▶ linear algebra *engine* in MATLAB, Python, R, Mathematica, ...
- ▶ high performance when optimized for a specific system

# BLAS

## BLAS level 1 routines (1970s)

- ▶ vector operations, e.g.,

$$x^T y, \quad \|x\|_2, \quad x \leftarrow \alpha x, \quad y \leftarrow \alpha x + y$$

- ▶ use  $O(n)$  operations for vectors of length  $n$

## BLAS level 2 routines (1980s)

- ▶ matrix-vector operations, e.g.,

$$y \leftarrow \alpha Ax + \beta y, \quad A \leftarrow \alpha xx^T + A, \quad x \leftarrow T^{-1}b, \quad T \text{ triangular}$$

- ▶ use  $O(mn)$  operations for matrices of size  $m \times n$

# BLAS

## BLAS level 3 routines (1980s)

- ▶ matrix-matrix routines, e.g.,

$$C \leftarrow \alpha AB + \beta C, \quad X \leftarrow T^{-1}B, \quad T \text{ triangular}$$

- ▶ use  $O(n^3)$  operations for matrices of size  $n \times n$

# What's in a name?

## BLAS naming scheme

XYYZZ

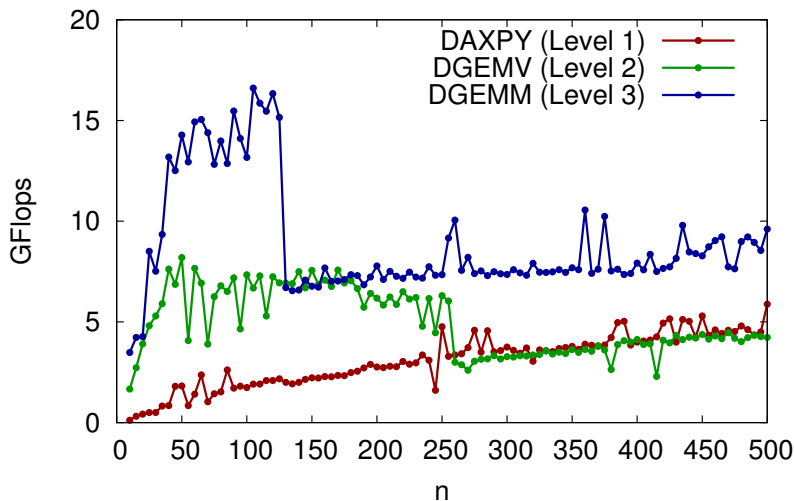
- ▶ First character X indicates data type (S, D, C, Z)
- ▶ BLAS level 1: letters YYZZ indicate mathematical operation
- ▶ BLAS level 2+3: letters YY indicate matrix type
- ▶ BLAS level 2+3: letters ZZ indicate mathematical operation

## Examples

- ▶ dscale — *double scale* ( $x \leftarrow \alpha x$ )
- ▶ saxpy — *single a x plus y* ( $y \leftarrow \alpha x + y$ )
- ▶ dgemv — *double general matrix-vector* ( $y \leftarrow \alpha Ax + \beta y$ )
- ▶ dtrsv — *double triangular solve vector* ( $x \leftarrow T^{-1}x$ )
- ▶ ssymm — *single symmetric matrix-matrix* ( $C \leftarrow \alpha SB + \beta C$ )

# BLAS performance

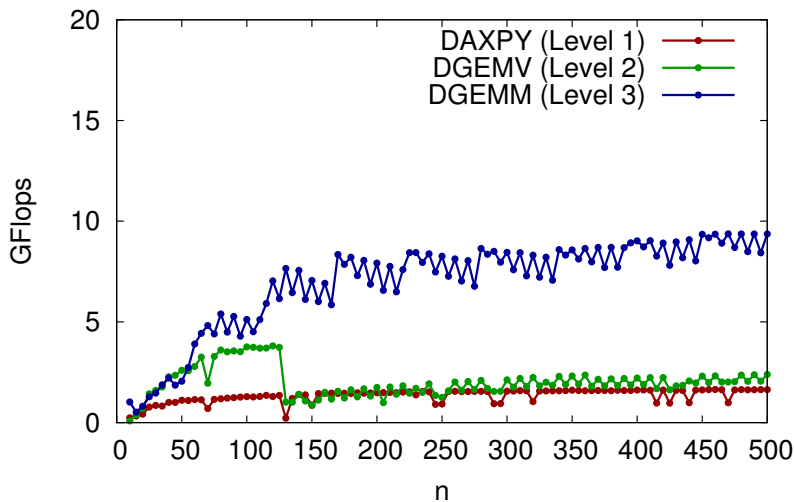
Intel Core i5 @ 1.8 GHz (L3 3MB, L2 256K, L1 32K)






# BLAS performance

Intel Xeon X5650 @ 1.6 GHz (L3 12MB, L2 256K, L1 32K)



# Memory hierarchy



Source: <http://goo.gl/yGtSDI>

## Possible bottlenecks

- ▶ CPU speed is a bottleneck — *CPU bound*
- ▶ cache size/speed is a bottleneck — *cache bound*
- ▶ memory size/speed is a bottleneck — *memory bound*
- ▶ disk or network speed is a bottleneck — *I/O bound*

# BLAS memory model

- ▶ vectors and matrices are contiguous arrays
- ▶ matrices are stored in column-major ordering
- ▶ *stride* refers to distance between *consecutive* elements
- ▶ *leading dimension* (LDA) refers to distance between columns

$$A = \begin{bmatrix} A_{11} & A_{12} & A_{13} & A_{14} & A_{15} \\ A_{21} & A_{22} & A_{23} & A_{24} & A_{25} \\ A_{31} & A_{32} & A_{33} & A_{34} & A_{35} \\ A_{41} & A_{42} & A_{43} & A_{44} & A_{45} \end{bmatrix}, \quad \begin{bmatrix} * & * & * & * & * \\ * & * & A_{23} & A_{24} & A_{25} \\ * & * & A_{33} & A_{34} & A_{35} \\ * & * & * & * & * \end{bmatrix}$$

- ▶  $i$ th column of  $A$  is a vector of length 4 with stride 1
- ▶  $i$ th row of  $A$  is a vector of length 5 with stride 4
- ▶  $(A_{11}, A_{22}, A_{33}, A_{44})$  is a vector of length 4 with stride 5
- ▶  $A$  is a matrix with 4 rows, 5 columns, stride 1, LDA 4
- ▶ *slice* (submatrix to the right) has 2 rows, 3 columns, stride 1, LDA 4

## Example: BLAS level 1

```
/* Prototype for BLAS dscal */
void dscal_(
    const int * n,           /* length of array */
    const double * a,        /* scalar a */
    double * x,              /* array x */
    const int * incx         /* array x, stride */
);

int main(void) {
    int i, incx, n;
    double a, x[5] = {2.0, 2.0, 2.0, 2.0, 2.0};

    /* Scale the vector x by 3.0 */
    n = 5; a = 3.0; incx = 1;
    dscal_(&n, &a, x, &incx);

    return 0;
}
```

# Makefile (DTU Unix system)

```
CC=gcc
SRCS=example_blas.c
EXECUTABLE=example_blas
INCLUDES=
CFLAGS=-c -Wall
LFLAGS=-L/usr/lib64/atlas    # directory that contains library
LIBS=-lf77blas              # name of BLAS library
OBJS=$(SRCS:.c=.o)

all: $(SRCS) $(EXECUTABLE)

$(EXECUTABLE): $(OBJS)
    $(CC) $(OBJS) $(LFLAGS) $(LIBS) -o $@

%.o: %.c
    $(CC) -c $(CFLAGS) $(INCLUDES) $< -o $@

clean:
    rm *.o $(EXECUTABLE)
```

## Example: CBLAS

```
#include <stdio.h>
#ifdef __MACH__ && defined(__APPLE__)
#include <Accelerate/Accelerate.h>
#else
#include <cbblas.h>
#endif

int main(void) {

    int i, incx, n;
    double a, x[5] = {2.0, 2.0, 2.0, 2.0, 2.0};

    /* Scale the vector x by 3.0 */
    n = 5; a = 3.0; incx = 1;
    cbblas_dscal(n, a, x, incx);

    return 0;
}
```

# Makefile (Windows)

```
CC=gcc
SRCS=example_cblas.c
EXECUTABLE=example_cblas.exe
INCLUDES=-I.                # directory that contains cblas.h
CFLAGS=-c -Wall
LFLAGS=-L.                  # directory that contains libraries
LIBS=-lopenblas            # name of BLAS & CBLAS libraries
OBJS=$(SRCS:.c=.o)

all: $(SRCS) $(EXECUTABLE)

$(EXECUTABLE): $(OBJS)
    $(CC) $(OBJS) $(LFLAGS) $(LIBS) -o $@

%.o: %.c
    $(CC) -c $(CFLAGS) $(INCLUDES) $< -o $@

clean:
    rm *.o $(EXECUTABLE)
```